



UPPSALA
UNIVERSITET

GPUs

Carl Nettelblad

2020-06-01



Welcome to the course!

- A course on how to scale machine learning models
 - And other computations.
- How do we scale from CPU to GPU?
 - Why does it work?
 - When won't it work?
 - How can we roll our own algorithms?
- How do we scale from *development* on single nodes to *production workflows* in a cloud?
- A project where you can try these techniques out yourself.

Who am I?

- Carl Nettelblad
 - PhD in Scientific Computing 2012
 - Main research in statistical methods for life science
 - Much implemented on CPUs, but interests include deep learning, and other processing, on GPUs
 - Supervised my first student project using CUDA in 2009
 - Currently associate professor at Department of IT, UU, as well as technical coordinator at UPPMAX
 - A past as a competitor in programming competitions, including the International Olympiad in Informatics, the ICPC world finals, and TopCoder Open
 - I love the intersection of theory, algorithms, and applications. I strongly believe the best science is made when you take a small step outside of run-of-the-mill approaches in all three.

Please!

- After the first break, I'll ask you all to just say some quick words about yourself.
- Who are you? What department/university? Are you taking this course as a way to broaden your views, or do you see specific connections to your research?

Interaction!

- Please interrupt!
- It is harder to “read the room” and see what needs clarification or not, when there is no room.
- You all have diverse experience and knowledge that is worth sharing.

What's a CPU?

- A von Neumann machine
- Memory
 - Instructions and data
- CPU
 - Read instruction from current instruction counter, increment counter
 - Instruction might request data from memory
 - Instruction might modify instruction counter
 - Repeat until turned off

What's a CPU?

- This is a fairytale.
- Very early CPUs worked something like this.
- One instruction at a time, instruction i has to complete before executing instruction $i + 1$.

Why is this not true?

Intel Skylake

- The CPU is:
 - Pipelined, several instructions are “in flight” at once
 - Read and decode instruction $i + 2$
 - Compute the result of instruction $i + 1$
 - Store the result of instruction i
 - Superscalar, dispatch several instruction per clock cycle
 - Several pipeline steps are “wide”, allowing multiple instructions to be handled at once
 - Out-of-order
 - If instruction i can not complete due to waiting for data, or the result of another instruction, start instruction $i + 1$ instead
 - Cached
 - Main memory is slow, store copies locally
 - All of these had happened by the year 2000!

14-
19

5

224

256
kb

But there's more...

- The vendors added more and more logic to keep the illusion of the von Neumann machine.
- All these steps still maintain the illusion of a single stream of instructions (a single thread) executing at once.
- This also includes tricks like *branch predictors* and *prefetchers*.
 - You need to know the result of an `if` to know what the next instruction is
 - You want to guess what indexes will be accessed in order to (ideally) bring those into the cache before they are accessed!
- Mispredictions and cache misses are expensive.
 - Seemingly innocent computer science concepts like virtual methods and linked lists can get expensive when realized on a modern architecture

What happened after 2000?

- Frequencies topped out at 5 GHz
- It got exceedingly hard to keep a single thread running fast
- Add more cores
 - Don't just add execution units within the core, but fully separate CPUs within the same socket
 - 4-8 common in desktops and high-end laptops now
 - Add simultaneous multi-threading, 2-4 threads on the *same* core
- Add more cache
 - Shared between cores, L3
- Add wider *vector* instructions
 - Process more 32-bit numbers at the same time
 - 4... 8... 16... 32
 - But even HPC codes have a hard time to use these

Single core CPU die shot



From wikichip

This is what makes CPU great!

- They are (comparatively) “easy to program” for the very reason that most of the silicon is spent on hiding problems and speeding up the simple von Neumann abstraction



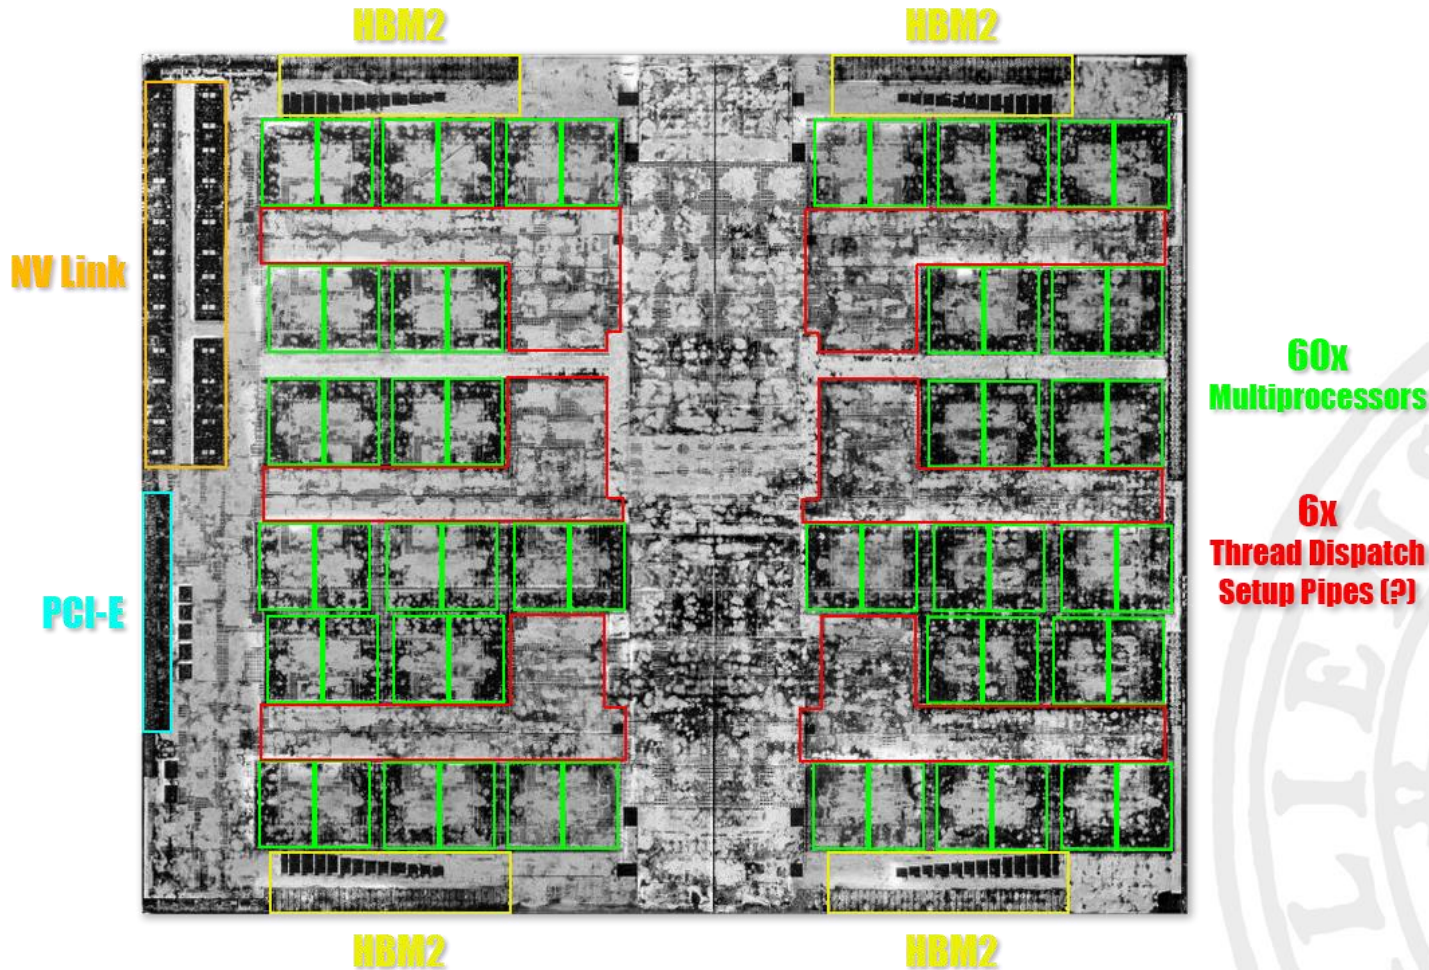
What's a GPU?

- Render pixels
 - In the year 2000, 1024x768 at 60 FPS
 - Done by drawing triangles
 - Compute the position of each triangle, compute the color of each pixel
- Extremely parallel
- Made more flexible over time, but the main principle stands
 - There is *not* a single main thread or a few of them
 - There are thousands of threads



UPPSALA
UNIVERSITET

GPU die shot



The GPU is limited

- Not out of order
- No branch prediction
- Very limited automatic prefetching
- Some caches
- Lots of simultaneous multithreading
 - When stalling (waiting for operation to complete), switch to another thread
- So, if you have 1000 “cores”, you can easily want to have 8000 threads
- **Does your algorithm have 8000 threads?**

Memory

- Normal system memory is slow (maybe 30 GB/s/CPU)
- A GPU can easily have 300 GB/s
- Lower latency as well
- Again, the CPU cores create this fairytale world
 - Out of order, branch prediction, prefetching, lots of cache space *per thread*
 - And much of caches local per core
 - You want to find independent tasks with independent data

Comparison

- A Rackham node
 - 20 cores
 - At most 40 simultaneous threads
 - 128 GB system memory
 - ~60 GB/s memory bandwidth
 - 5120 kB core-specific cache (only counting L2)
 - 250+ W just for CPUs

Comparison

- Tesla T4
 - A lightweight current-generation GPU
 - A gaming card redressed for servers
 - 2560 cores
 - At most 40,960 simultaneous threads
 - 16 GB video memory
 - ~300 GB/s memory bandwidth
 - 3840 kB SM-specific memory/cache in total
 - Shared by 64 cores
 - 70 W per GPU

Resources per thread

Resource/thread	Rackham CPU	Snowy GPU
Cores	0.5	0.0625
Memory	3250 MB	0.4 MB
Memory bandwidth	1500 MB/s	7.5 MB/s
Cache	131 072 bytes	96 bytes
Scalar SP FLOPS (base)	2200 GFLOPS	73 GFLOPS
Scalar SP FLOPS (boost)	3400 GFLOPS	200 GFLOPS
Vector SP FLOPS (base)	14 400 GFLOPS	73 GFLOPS
Vector SP FLOPS (boost)	24 800 GFLOPS	200 GFLOPS

Math

- Originally, GPUs focused on rendering RGB pixels, 3*8-bit colors
- Later generalized to floating point
 - FP good for all 3D transforms
- These days, they can do proper integer math just as fast
- Text is just integers
- Memory addresses are also integers

Summary

- An extremely parallel computational unit
 - But fully programmable
- For CPU
 - Keep threads separate
 - Thread count ~ core count
 - Communication is expensive
 - Vector instructions for high performance
- For GPU
 - Each thread is weak
 - Thread count >> core count
 - Keep multiple threads working in concert
 - Groups of threads mainly replace vector instructions



Nvidia's take on this

TRADITIONAL ASSUMPTIONS

GPUS ARE GOOD FOR:

- Floats, short floats, and doubles.
- Arrays (may be multi-dimensional).
- Coalesced memory access.
- Lock-free algorithms.
- Arithmetic.
- Compute-bound workloads.

GPUS ARE BAD FOR:

- Strings.
- Complex data structures.
- Random memory walks.
- Starvation-free algorithms (spinlocks).
- Control flow.
- Memory-BW/latency-bound workloads.

REALITY

NVIDIA GPUS ACCELERATE CODE DOING ALL OF THE ABOVE

My take on Nvidia

- The left-hand box is *easier* to make fast
- But those algorithms also shine on CPUs
 - GPUs excel at dense matrix multiplications and multi-dimensional Fourier transforms
 - These algorithms have repeated passes over the same data
 - But CPU vector instructions also do reasonably well
- The single most important thing is extreme parallelism
 - *Relative to other resources*
 - The relative benefit to CPU can be the greatest when the algorithm is hard to coax into vector instructions (manually or using compiler)
 - Even if resource utilization of GPU gets abysmal

Please!

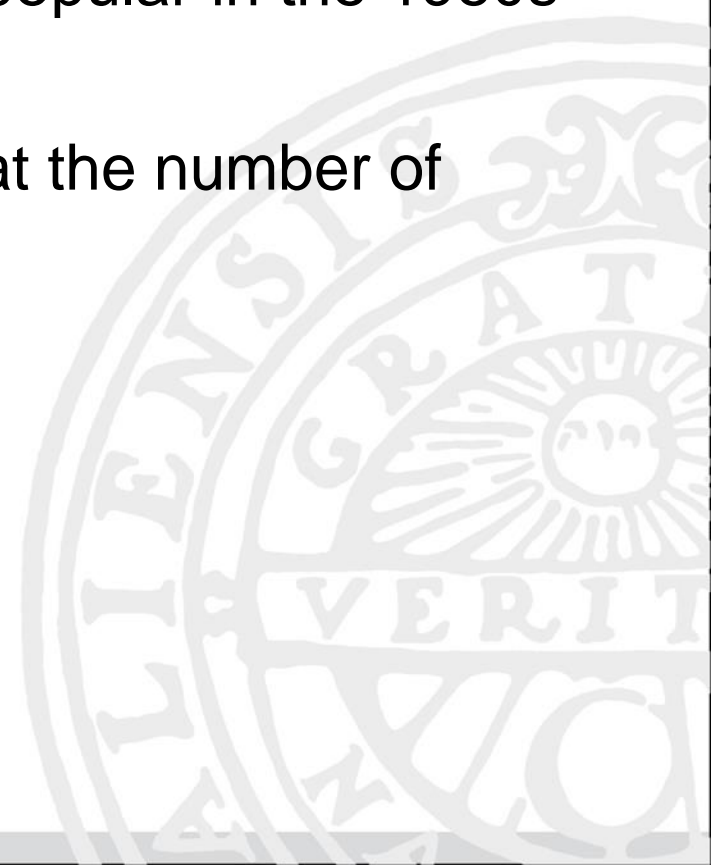
- After the first break, I'll ask you all to just say some quick words about yourself.
- Who are you? What department/university? Are you taking this course as a way to broaden your views, or do you see specific connections to your research?

Deep learning

- We focus more on the computational machinery in this course than actual models
- Still relevant to have some familiarity
- I know several of you actively work with deep models
 - I welcome discussion and clarifications

Deep learning


- Artificial neural networks have existed since the 1960s
 - “Multi-layer perceptrons” were popular in the 1980s
- Deep learning refers to the fact that the number of layers has increased significantly





Some layers

- A simple two-layer network, inputs x , outputs y
- Inputs are multiplied by weights
 - Summed
- Bias added
- Activation function applied
- Result from activation functions forms output, fed to the next layer



Graph

- The network is a directed (acyclic) graph
 - This is called a feed-forward network
 - General, biological, neural circuits do not need to behave like this
 - But in fact, some important biological structures *do*
- The edges indicate the flow of data

Backpropagation

- Define your desired outputs (y in our example)
- Define how to compare outputs against the desired outputs (objective function/loss function)
 - Mean square error is a simple example
- How can we minimize the loss?
 - Compute the gradient
 - If you know the activation and the gradient for the output step, you can also compute the gradient for any variable and activation function backwards
- Take a small step in the direction of lower loss value for every single optimizable variable in the network

Mini-batches

- The training dataset is frequently very large
- To do backpropagation, you need to store the activation value for every single node in the network, for every input
 - This consumes space and bandwidth
- Especially on GPUs, it is therefore crucial to instead train in batches
 - More samples in a batch might allow greater compute parallelism
 - More samples in a batch can exhaust memory and reduce cache performance

Optimization

- Just tracking the gradient seems simple
 - But it's extremely slow
 - Small steps (low *learning rate*) not to overshoot
 - Function not convex
 - Moving towards the top of a hill won't bring you to Mount Everest

Optimization

- Different batch configurations and optimization algorithms try to avoid this
 - If you use different random subsets for each batch, hopefully the stable gradient signals indicate “true” information
 - Optimization algorithms employ momentum, accumulating a trend in the gradient over many iterations
 - Adam (Adaptive momentum) a very popular choice

Optimization

- Attempts to avoid rough model landscapes of many local minima and *overfitting* to only training set, not general behavior
 - Dropout
 - Randomly removing some activations at each training instance
 - Try to make sure that no single value controls the whole thing
 - Batch normalization
 - Reshaping activations to approximate a standard normal distribution
 - Based on the set of activations within the mini-batch

Convolution

- Our example was a fully connected network
 - Also called “dense”
 - Each node in each layer is connected to every node in the next layer by a weight
- A fully connected network will have a very high number of weights
 - Large model
 - I.e. slow
 - Hard to train
 - Lots of overfitting/local minima
 - Limited ability to generalize

Ask your neighbor

- In a time series, or an image (or a combination, like a video), the absolute location is not crucial
 - Rather, compare to the adjacent samples
- A *convolutional* neural network is *locally connected* with *shared weights*
 - Apply a small (e.g. 3x3, 5x5) stencil to each pixel
 - Apply a shared set of weights to the pixels from the previous layer
 - If this is made deep, information can move far further than just one stencil width
 - Also reduction operations (averaging, max pooling) to reduce dimensionality

What does this mean for performance?

- Fully connected
 - Each of N_1 activations is used N_2 times
 - Lots of reuse of data!
 - But reuse is global
- Convolutional
 - Each of N_1 activations is used *stencil size* times
 - These are close together!
 - You can imagine one GPU thread computing the activation from each new pixel
 - Adjacent GPU threads will access mostly the same data

Adversarial networks

- You can get a low mean-square error for reconstructing an image, and still get a result that is obviously and clearly wrong if you just look at it
- One option is to train multiple networks
 - Generator – the “main” network
 - Generate a face from a random seed vector
 - Discriminator – judge the result of the main network
 - Answer: Is this a natural face?
- Discriminator is trained on true examples and generated ones
- Generator trained on signal from discriminator
- Adversarial since they are trained for opposite goals
 - Discriminator adapts to identify typical generator errors
 - Generator adapts to fool discriminator

What does this mean for performance?

- Lots of computations going on
- More data locally on the GPU
- Rather than (only) feeding in training data from an outside source, data is generated and used locally
- Adversarial networks have formed the basis for many of the most innovative developments during the last few years
 - Defining a good error metric used to be a hard challenge

Pre-lab 1

- What results did you get?
- Can you relate those to what we have talked about?
(What kind of model did we use?)





TensorFlow

- Library from Google
- Based on creating the computational graph through a syntax rather similar to using numpy in Python
 - But not identical
- If you have access to the full graph, you can optimize memory transfers and evaluation order
- Lots of constructs specifically adapted to neural networks
 - Convolution operations, loss functions, optimizers
- But also a general evaluator of any computational graph

What is a Tensor?

- In practice, a TensorFlow tensor is an object similar to a numpy ndarray.
- It has a dimensionality and a type.
- Even in numpy, not every array-like object is backed by an actual chunk of memory.
 - If you slice or transpose an array, that only creates a view inside another array.
- In TensorFlow, many operations are just that, if you write $C = A * B$, C represents the operation of multiplying A and B , in their current state.
 - That operation might also be executed.

Constants, variables

- You want to insert a value (scalar or an array) into TensorFlow
 - `tf.constant`, e.g.
 - `tf.constant(someData, dtype=tf.float32)`
- If you just want to create a large tensor of identical values, `tf.fill` is more efficient
 - Accepting a scalar and a shape
- You can explicitly create optimizable variables with `tf.Variable`
 - These are end targets for gradient computations and optimization algorithms
 - The actual content of a model

Operations

- Just like in numpy, operations are elementwise per default.
- $A * B$, $A + B$, A / B all act elementwise
- $A @ B$ for matrix multiplication (syntactic sugar since Python 3.5, implies `tf.matmul`)
- `tf.stack` combines several tensors of rank R into one of rank $R + 1$
- `tf.concat` concatenates several tensors of rank R into a new tensor of rank R
- `tf.where` takes a boolean tensor choosing elements from the arguments, e.g. maximum by `tf.where(A > B, A, B)`
 - `tf.math.maximum(A, B)` is more convenient and probably more efficient
 - `tf.math.reduce_max(A)` instead computes the maximum within a tensor

General arguments

- Almost all TensorFlow functions accept a name argument. TensorFlow itself doesn't "see" your variable names, so error messages can sometimes refer to layer names.
- Many functions can work on the full tensor, or just along some axis. The operation mode can be changed using the axis argument.
 - `tf.reduce_sum` has a default axis of None, summing over all axes.
 - `tf.concat` and `tf.stack` have defaults of 0, indicating that the common/new axis should be the first one.
- Negative indices are allowed, just like in ordinary Python.

Eager execution

- You might read a lot of stuff on how one first creates a graph and then asks TensorFlow to run to get the value of some specific tensor
 - This was the only viable usage mode in TensorFlow 1
 - Create the graph, then run it
 - TensorFlow 2 supports eager execution, where the operations you do are also evaluated
 - This is *much* easier to debug... but it removes some optimization opportunities

@tf.function

- As stated, TensorFlow will run in eager mode by default
- You can decorate a function by `@tf.function`
- This will make TensorFlow analyze the full function and try to express it as TensorFlow graphs
- This gives room for optimization
- Sometimes even for loops and other “expensive” things in Python can be pushed into a compact graph that is executed all on the GPU
- Enclose well-contained logic in functions (good practice anyway), and tag them as `@tf.function`, unless you want to debug them

Pitfalls

- A tensor can stay on the GPU
- TensorFlow can be very helpful in converting to and from numpy arrays
 - If you do that, the information will be transferred CPU \leftrightarrow GPU
 - All gradient information will be lost if you convert from a tensor to numpy, do something, and then transform back
 - From TensorFlow's point of view, those are two unrelated constant tensors
- Even just doing an if on some value of a tensor outside of a `@tf.function` forces the full evaluation of the tensor and the transfer of that data from GPU to CPU
- It's so easy to do some things that it's not clear that some things are expensive

Broadcasting

- Like numpy, TensorFlow supports broadcasting
 - If a tensor is size 1 in some dimension, that can implicitly be interpreted to match any other size for many operations
- Also creates risk for hard-to-understand errors
- Multiplying a shape $(1,N)$ vector with a $(N,1)$ vector creates a (N,N) matrix
- `tf.broadcast_to`, `tf.reshape` and `tf.expand_dims` can be useful when you need a bit more control
- Broadcasting is far more efficient than actually creating the corresponding tensor with repeated elements
- The same holds for `tf.tile`, for repeating a tensor multiple times

Keras

- In the prelab, you saw an example of Keras.
- Keras is a higher level abstraction for building neural networks.
 - In theory, Keras can be used with several backends (not only TensorFlow), but it can also just be used as a convenience layer
- Sometimes confusing to find Keras and TensorFlow features to do the same thing.

Auto-differentiation

- Backpropagation would be terrible to use if you would implement it manually
- TensorFlow provides automatic computation of gradients
 - Automatic for variables, but you can compute the gradient for a constant as well

Consequences of differentiation

- (Almost) *any* TensorFlow operation is differentiable.
 - But the gradient of e.g. a max operation is only related to the maximum value. Even if the lower value is only a factor $1-1e-5$ smaller, it's ignored.
 - This and other discrete operations is something you want to avoid. Rather, one tries to have “soft” operations, like computing a norm, or taking the sum of the logged exponentials.

Code example

```
x = tf.constant(3.0)
with tf.GradientTape(persistent=True) as g:
    g.watch(x)
    y = x * x
    z = y * y
dz_dx = g.gradient(z, x) # 108.0 (4*x^3 at x = 3)
dy_dx = g.gradient(y, x) # 6.0
del g # Drop the reference to the tape
```

Other frameworks

- PyTorch
 - An early main difference was that PyTorch was more like the eager mode in TensorFlow 2
 - Somewhat easier to use for distributed training over multiple GPUs
 - Almost a matter of taste these days
 - Many models can be found for TensorFlow as well as PyTorch
 - Sometimes higher performance in benchmarks
- MXNet
 - More advanced and varied GPU support
 - Frequently higher performance in benchmarks
 - The basis for Minpy, that we will try tomorrow
 - Not as widespread in research or industrial usage, harder to find good online resources

Lab 1

- Game of Life, created by John Horton Conway in 1970
 - Conway passed away in Covid-19 earlier this year
- Discrete Cartesian grid of cells
 - Each grid cell has 8 neighbors
 - Each cell can be alive or dead



Standard rules

- Any live cell with two or three neighbors survive
 - All other live cells die
- Any dead cell with three live neighbors gets alive
 - All other dead cells stay dead
- All rules evaluated based on number of live/dead neighbors in previous timestep

Naïve algorithm

- Compute number of live neighbors to each cell on a fixed grid
 - Possibly with extra term for current cell alive/dead
- Apply rule based on live/dead status and number of live neighbors
- This can be represented as a convolution
- Repeat for N generations

Hashlife

- Just for reference, the most advanced algorithm uses the fact that Game of Life is fully deterministic
 - If you find that a certain patch of the universe is identical to some patch you've seen before, you can predict the evolution of that patch
 - With some caveats of course
- We will not look into this



Example

- [https://en.wikipedia.org/wiki/Breeder_\(cellular_automaton\)#/media/File:Conways_game_of_life_breeder_animation.gif](https://en.wikipedia.org/wiki/Breeder_(cellular_automaton)#/media/File:Conways_game_of_life_breeder_animation.gif)
- https://upload.wikimedia.org/wikipedia/commons/e/e6/Conways_game_of_life_breeder_animation.gif

Q&A

- See you at 1.15 PM
 - You don't have to stay in the Zoom room all the time
 - I'll try to make sure that I give notice on Slack if we are about to discuss something of interest to a lot of people
- Feel free to start with the lab before that
- Note that the reservation is only during 1.00-5.00, so if you want to work on it outside that time, do not use the --reservation flag